

How to use dates & times with pandas

MANIPULATING TIME SERIES DATA IN PYTHON



Stefan Jansen

Founder & Lead Data Scientist at Applied
Artificial Intelligence

Date & time series functionality

- At the root: data types for date & time information
 - Objects for points in time and periods
 - Attributes & methods reflect time-related details
- Sequences of dates & periods:
 - Series or DataFrame columns
 - Index: convert object into Time Series
- Many Series/DataFrame methods rely on time information in the index to provide time-series functionality

Basic building block: pd.Timestamp

```
import pandas as pd # assumed imported going forward
from datetime import datetime # To manually create dates

time_stamp = pd.Timestamp(datetime(2017, 1, 1))

pd.Timestamp('2017-01-01') == time_stamp
```

Timestamp is the pandas equivalent of python's Datetime and is interchangeable with it in most cases.

```
True # Understands dates as strings
```

```
time_stamp # type: pandas.tslib.Timestamp
```

If you display the timestamp, you'll notice that the time has been automatically set to midnight.

```
Timestamp('2017-01-01 00:00:00')
```

Basic building block: pd.Timestamp

- Timestamp object has many attributes to store time-specific information

```
time_stamp.year
```

```
2017
```

```
time_stamp.weekday_name
```

```
'Sunday'
```

More building blocks: pd.Period & freq

```
period = pd.Period('2017-01')  
period # default: month-end
```

```
Period('2017-01', 'M')
```

```
period.asfreq('D') # convert to daily
```

```
Period('2017-01-31', 'D')
```

```
period.to_timestamp().to_period('M')
```

```
Period('2017-01', 'M')
```

- Period object has freq attribute to store frequency info It has a method to convert between frequencies.

- Convert pd.Period() to pd.Timestamp() and back

More building blocks: pd.Period & freq

```
period + 2
```

```
Period('2017-03', 'M')
```

```
pd.Timestamp('2017-01-31', 'M') + 1
```

```
Timestamp('2017-02-28 00:00:00', freq='M')
```

- Frequency info enables basic date arithmetic

Sequences of dates & times

- `pd.date_range` : `start` , `end` , `periods` , `freq` The default is daily frequency.

```
index = pd.date_range(start='2017-1-1', periods=12, freq='M')
```

```
index
```

```
DatetimeIndex(['2017-01-31', '2017-02-28', '2017-03-31', ...,  
              '2017-09-30', '2017-10-31', '2017-11-30', '2017-12-31'],  
              dtype='datetime64[ns]', freq='M')
```

- `pd.DatetimeIndex` : sequence of Timestamp objects with frequency info

Sequences of dates & times

```
index[0]
```

```
Timestamp('2017-01-31 00:00:00', freq='M')
```

```
index.to_period()
```

```
PeriodIndex(['2017-01', '2017-02', '2017-03', '2017-04', ...,  
            '2017-11', '2017-12'], dtype='period[M]', freq='M')
```


Create a time series: `pd.DatetimeIndex`

```
pd.DataFrame({'data': index}).info()
```

```
RangeIndex: 12 entries, 0 to 11  
Data columns (total 1 columns):  
data      12 non-null datetime64[ns]  
dtypes: datetime64[ns](1)
```

Create a time series: `pd.DatetimeIndex`

- `np.random.random` :
 - Random numbers: `[0, 1]`
 - 12 rows, 2 columns

```
data = np.random.random((size=12,2))  
pd.DataFrame(data=data, index=index).info()
```

```
DatetimeIndex: 12 entries, 2017-01-31 to 2017-12-31  
Freq: M  
Data columns (total 2 columns):  
0      12 non-null float64  
1      12 non-null float64  
dtypes: float64(2)
```

Frequency aliases & time info

There are many frequency aliases besides 'M' and 'D':

Period	Alias
Hour	H
Day	D
Week	W
Month	M
Quarter	Q
Year	A

These may be further differentiated by beginning/end of period, or business-specific definition

You can also access these `pd.Timestamp()` attributes:

attribute
.second, .minute, .hour,
.day, .month, .quarter, .year
.weekday
dayofweek
.weekofyear
.dayofyear

Let's practice!


MANIPULATING TIME SERIES DATA IN PYTHON

Indexing & resampling time series

MANIPULATING TIME SERIES DATA IN PYTHON

Stefan Jansen

Founder & Lead Data Scientist at Applied Artificial Intelligence



Time series transformation

Basic time series transformations include:

- Parsing string dates and convert to `datetime64`
- Selecting & slicing for specific subperiods
- Setting & changing `DateTimeIndex` frequency
 - Upsampling vs Downsampling

Getting GOOG stock prices

```
google = pd.read_csv('google.csv') # import pandas as pd
google.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 504 entries, 0 to 503
Data columns (total 2 columns):
date      504 non-null object
price     504 non-null float64
dtypes: float64(1), object(1)
```

```
google.head()
```

```
   date    price
0 2015-01-02  524.81
1 2015-01-05  513.87
2 2015-01-06  501.96
3 2015-01-07  501.10
4 2015-01-08  502.68
```

Converting string dates to datetime64

- `pd.to_datetime()` :
 - Parse date string
 - Convert to `datetime64`

```
google.date = pd.to_datetime(google.date)
google.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 504 entries, 0 to 503
Data columns (total 2 columns):
date      504 non-null datetime64[ns]
price     504 non-null float64
dtypes: datetime64[ns](1), float64(1)
```


Converting string dates to datetime64

- `.set_index()` :
 - Date into index
 - `inplace` :
 - `don't create copy`

```
google.set_index('date', inplace=True)
google.info()
```


```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 504 entries, 2015-01-02 to 2016-12-30
Data columns (total 1 columns):
price      504 non-null float64
dtypes: float64(1)
```

Plotting the Google stock time series

```
google.price.plot(title='Google Stock Price')
```

```
plt.tight_layout(); plt.show()
```

With a `DateTime` Index, pandas automatically creates reasonably spaced date labels for the x axis.



Partial string indexing

- Selecting/indexing using strings that parse to dates

```
google['2015'].info() # Pass string for part of date
```

```
DatetimeIndex: 252 entries, 2015-01-02 to 2015-12-31  
Data columns (total 1 columns):  
price      252 non-null float64  
dtypes: float64(1)
```

```
google['2015-3': '2016-2'].info() # Slice includes last month  
includes the end date
```

```
DatetimeIndex: 252 entries, 2015-03-02 to 2016-02-29  
Data columns (total 1 columns):  
price      252 non-null float64  
dtypes: float64(1)  
memory usage: 3.9 KB
```

Partial string indexing

```
google.loc['2016-6-1', 'price'] # Use full date with .loc[]
```

```
734.15
```

.asfreq(): set frequency

- `.asfreq('D')` :
 - Convert `DatetimeIndex` to calendar day frequency

```
google.asfreq('D').info() # set calendar day frequency
```

```
DatetimeIndex: 729 entries, 2015-01-02 to 2016-12-30
Freq: D
Data columns (total 1 columns):
price      504 non-null float64
dtypes: float64(1)
```

.asfreq(): set frequency

- Upsampling:
 - Higher frequency implies new dates => missing data

```
google.asfreq('D').head()
```

As a result, the `DatetimeIndex` now contains many dates where stock wasn't bought or sold.

	price
date	
2015-01-02	524.81
2015-01-03	NaN
2015-01-04	NaN
2015-01-05	513.87
2015-01-06	501.96

.asfreq(): reset frequency

- `.asfreq('B')` :
 - Convert `DatetimeIndex` to `business day frequency`

```
google = google.asfreq('B') # Change to calendar day frequency
google.info()
```

```
DatetimeIndex: 521 entries, 2015-01-02 to 2016-12-30
Freq: B
Data columns (total 1 columns):
price      504 non-null float64
dtypes: float64(1)
```

.asfreq(): reset frequency

```
google[google.price.isnull()] # Select missing 'price' values
```

	price
date	
2015-01-19	NaN
2015-02-16	NaN
...	
2016-11-24	NaN
2016-12-26	NaN

- Business days that were not trading days

Let's practice!


MANIPULATING TIME SERIES DATA IN PYTHON

Lags, changes, and returns for stock price series

MANIPULATING TIME SERIES DATA IN PYTHON

Stefan Jansen

Founder & Lead Data Scientist at Applied
Artificial Intelligence



Basic time series calculations

- Typical Time Series manipulations include:
 - Shift or lag values back or forward back in time
 - Get the difference in value for a given time period
 - Compute the percent change over any number of periods
- `pandas` built-in methods rely on `pd.DatetimeIndex`

Getting GOOG stock prices

- Let `pd.read_csv()` do the parsing for you!

```
google = pd.read_csv('google.csv', parse_dates=['date'], index_col='date')
```

```
google.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 504 entries, 2015-01-02 to 2016-12-30
Data columns (total 1 columns):
price      504 non-null float64
dtypes: float64(1)
```

Getting GOOG stock prices

```
google.head()
```

```
      price  
date  
2015-01-02  524.81  
2015-01-05  513.87  
2015-01-06  501.96  
2015-01-07  501.10  
2015-01-08  502.68
```

.shift(): Moving data between past & future

- `.shift()` :
 - defaults to `periods=1`
 - 1 period into future

```
google['shifted'] = google.price.shift() # default: periods=1
google.head(3)
```

	price	shifted
date		
2015-01-02	542.81	NaN
2015-01-05	513.87	542.81
2015-01-06	501.96	513.87

.shift(): Moving data between past & future

- `.shift(periods=-1)` :
 - lagged data
 - 1 period back in time

```
google['lagged'] = google.price.shift(periods=-1)
google[['price', 'lagged', 'shifted']].tail(3)
```

	price	lagged	shifted
date			
2016-12-28	785.05	782.79	791.55
2016-12-29	782.79	771.82	785.05
2016-12-30	771.82	NaN	782.79

Calculate one-period percent change

- x_t / x_{t-1}

```
google['change'] = google.price.div(google.shifted)
google[['price', 'shifted', 'change']].head(3)
```

	price	shifted	change
Date			
2017-01-03	786.14	NaN	NaN
2017-01-04	786.90	786.14	1.000967
2017-01-05	794.02	786.90	1.009048

Calculate one-period percent change

```
google['return'] = google.change.sub(1).mul(100)
google[['price', 'shifted', 'change', 'return']].head(3)
```

	price	shifted	change	return
date				
2015-01-02	524.81	NaN	NaN	NaN
2015-01-05	513.87	524.81	0.98	-2.08
2015-01-06	501.96	513.87	0.98	-2.32

.diff(): built-in time-series change

- Difference in value for two adjacent periods
- $x_t - x_{t-1}$

```
google['diff'] = google.price.diff()  
google[['price', 'diff']].head(3)
```

	price	diff
date		
2015-01-02	524.81	NaN
2015-01-05	513.87	-10.94
2015-01-06	501.96	-11.91

.pct_change(): built-in time-series % change

- Percent change for two adjacent periods
- $\frac{x_t}{x_{t-1}}$

```
google['pct_change'] = google.price.pct_change().mul(100)
google[['price', 'return', 'pct_change']].head(3)
```

	price	return	pct_change
date			
2015-01-02	524.81	NaN	NaN
2015-01-05	513.87	-2.08	-2.08
2015-01-06	501.96	-2.32	-2.32

Looking ahead: Get multi-period returns

```
google['return_3d'] = google.price.pct_change(periods=3).mul(100)
google[['price', 'return_3d']].head()
```

	price	return_3d
date		
2015-01-02	524.81	NaN
2015-01-05	513.87	NaN
2015-01-06	501.96	NaN
2015-01-07	501.10	-4.517825
2015-01-08	502.68	-2.177594

- Percent change for two periods, 3 trading days apart

Let's practice!

MANIPULATING TIME SERIES DATA IN PYTHON