

# Indexing time series

PANDAS FOUNDATIONS



**Dhavide Aruliah**

Director of Training, Anaconda

# Using pandas to read datetime objects

- read\_csv() function
  - Can read strings into datetime objects
  - Need to specify 'parse\_dates=True'
- ISO 8601 format
  - yyyy-mm-dd hh:mm:ss

# Product sales CSV

|   | Date                | Company         | Product  | Units |
|---|---------------------|-----------------|----------|-------|
| 0 | 2015-02-02 08:30:00 | Hooli           | Software | 3     |
| 1 | 2015-02-02 21:00:00 | Mediacore       | Hardware | 9     |
| 2 | 2015-02-03 14:00:00 | Initech         | Software | 13    |
| 3 | 2015-02-04 15:30:00 | Streeplex       | Software | 13    |
| 4 | 2015-02-04 22:00:00 | Acme Coporation | Hardware | 14    |

# Parse dates

```
import pandas as pd
sales = pd.read_csv('sales-feb-2015.csv',
                    parse_dates=True, index_col= 'Date')
# parse all compatible columns as datetime objects
```

# Parse dates

```
sales.head()
```

| Date                | Company         | Product  | Units |
|---------------------|-----------------|----------|-------|
| 2015-02-02 08:30:00 | Hooli           | Software | 3     |
| 2015-02-02 21:00:00 | Mediacore       | Hardware | 9     |
| 2015-02-03 14:00:00 | Initech         | Software | 13    |
| 2015-02-04 15:30:00 | Streeplex       | Software | 13    |
| 2015-02-04 22:00:00 | Acme Coporation | Hardware | 14    |

# Parse dates

```
sales.info()
```

```
DatetimeIndex: 19 entries, 2015-02-02 08:30:00 to 2015-02-26 09:00:00  
Data columns (total 3 columns):  
Company      19 non-null object  
Product      19 non-null object  
Units        19 non-null int64  
dtypes: int64(1), object(2)  
memory usage: 608.0+ bytes
```

# Selecting single datetime

```
sales.loc['2015-02-19 11:00:00', 'Company']
```

```
'Mediacore'
```

# Selecting whole day

```
sales.loc[ '2015-2-5' ]
```

| Date                | Company         | Product  | Units |
|---------------------|-----------------|----------|-------|
| 2015-02-05 02:00:00 | Acme Coporation | Software | 19    |
| 2015-02-05 22:00:00 | Hooli           | Service  | 10    |



# Partial datetime string selection

- Alternative formats:
  - `sales.loc['February 5, 2015']`
  - `sales.loc['2015-Feb-5']`
- Whole month: `sales.loc['2015-2']`
- Whole year: `sales.loc['2015']`

# Selecting whole month

```
sales.loc['2015-2']
```

| Date                | Company         | Product  | Units |
|---------------------|-----------------|----------|-------|
| 2015-02-02 08:30:00 | Hooli           | Software | 3     |
| 2015-02-02 21:00:00 | Mediacore       | Hardware | 9     |
| 2015-02-03 14:00:00 | Initech         | Software | 13    |
| 2015-02-04 15:30:00 | Streeplex       | Software | 13    |
| 2015-02-04 22:00:00 | Acme Coporation | Hardware | 14    |
| 2015-02-05 02:00:00 | Acme Coporation | Software | 19    |
| 2015-02-05 22:00:00 | Hooli           | Service  | 10    |
| 2015-02-07 23:00:00 | Acme Coporation | Hardware | 1     |
| 2015-02-09 09:00:00 | Streeplex       | Service  | 19    |
| 2015-02-09 13:00:00 | Mediacore       | Software | 7     |
| 2015-02-11 20:00:00 | Initech         | Software | 7     |
| 2015-02-11 23:00:00 | Hooli           | Software | 4     |
| 2015-02-16 12:00:00 | Hooli           | Software | 10    |
| 2015-02-19 11:00:00 | Mediacore       | Hardware | 16    |
| ...                 |                 |          |       |

# Slicing using dates/times

```
sales.loc[ '2015-2-16' : '2015-2-20' ]
```

| Date                | Company   | Product  | Units |
|---------------------|-----------|----------|-------|
| 2015-02-16 12:00:00 | Hooli     | Software | 10    |
| 2015-02-19 11:00:00 | Mediacore | Hardware | 16    |
| 2015-02-19 16:00:00 | Mediacore | Service  | 10    |

# Convert strings to datetime

```
evening_2_11 = pd.to_datetime(['2015-2-11 20:00',  
                               '2015-2-11 21:00', '2015-2-11 22:00', '2015-2-11 23:00'])  
evening_2_11
```

```
DatetimeIndex(['2015-02-11 20:00:00', '2015-02-11 21:00:00',  
               '2015-02-11 22:00:00', '2015-02-11 23:00:00'],  
              dtype='datetime64[ns]', freq=None)
```

# Reindexing DataFrame

```
sales.reindex(evening_2_11)
```

Reindexing involves providing a new index and matching data as required. Here, the reindex method returns a new DataFrame with the four rows corresponding to the times in evening\_2\_11.

|            |          | Company | Product  | Units |
|------------|----------|---------|----------|-------|
| 2015-02-11 | 20:00:00 | Initech | Software | 7.0   |
| 2015-02-11 | 21:00:00 | NaN     | NaN      | NaN   |
| 2015-02-11 | 22:00:00 | NaN     | NaN      | NaN   |
| 2015-02-11 | 23:00:00 | Hooli   | Software | 4.0   |

# Filling missing values

```
sales.reindex(evening_2_11, method='ffill')
```

|            |          | Company | Product  | Units |
|------------|----------|---------|----------|-------|
| 2015-02-11 | 20:00:00 | Initech | Software | 7     |
| 2015-02-11 | 21:00:00 | Initech | Software | 7     |
| 2015-02-11 | 22:00:00 | Initech | Software | 7     |
| 2015-02-11 | 23:00:00 | Hooli   | Software | 4     |

```
sales.reindex(evening_2_11, method='bfill')
```

|            |          | Company | Product  | Units |
|------------|----------|---------|----------|-------|
| 2015-02-11 | 20:00:00 | Initech | Software | 7     |
| 2015-02-11 | 21:00:00 | Hooli   | Software | 4     |
| 2015-02-11 | 22:00:00 | Hooli   | Software | 4     |
| 2015-02-11 | 23:00:00 | Hooli   | Software | 4     |

# Let's practice!

PANDAS FOUNDATIONS

# Resampling time series data

PANDAS FOUNDATIONS



**Dhavide Aruliah**

Director of Training, Anaconda



# Sales data

```
import pandas as pd
sales = pd.read_csv('sales-feb-2015.csv',
                    parse_dates=True, index_col='Date')
sales.head()
```

| Date                | Company         | Product  | Units |
|---------------------|-----------------|----------|-------|
| 2015-02-02 08:30:00 | Hooli           | Software | 3     |
| 2015-02-02 21:00:00 | Mediacore       | Hardware | 9     |
| 2015-02-03 14:00:00 | Initech         | Software | 13    |
| 2015-02-04 15:30:00 | Streeplex       | Software | 13    |
| 2015-02-04 22:00:00 | Acme Coporation | Hardware | 14    |

# Resampling

- Statistical methods over different time intervals
  - `mean()`, `sum()`, `count()`, etc.
- Downsampling
  - reduce datetime rows to slower frequency
- Upsampling
  - increase datetime rows to faster frequency

# Aggregating means

```
daily_mean = sales.resample('D').mean()  
daily_mean
```

| Date       | Units |
|------------|-------|
| 2015-02-02 | 6.0   |
| 2015-02-03 | 13.0  |
| 2015-02-04 | 13.5  |
| 2015-02-05 | 14.5  |
| 2015-02-06 | NaN   |
| 2015-02-07 | 1.0   |
| 2015-02-08 | NaN   |
| 2015-02-09 | 13.0  |
| 2015-02-10 | NaN   |
| 2015-02-11 | 5.5   |
| 2015-02-12 | NaN   |
| 2015-02-13 | NaN   |
| 2015-02-14 | NaN   |

# Verifying

```
print(daily_mean.loc['2015-2-2'])
```

```
Units      6.0  
Name: 2015-02-02 00:00:00, dtype: float64
```

```
print(sales.loc['2015-2-2', 'Units'])
```

```
Date  
2015-02-02 08:30:00      3  
2015-02-02 21:00:00      9  
Name: Units, dtype: int64
```

```
sales.loc['2015-2-2', 'Units'].mean()
```

```
6.0
```

# Method chaining

```
sales.resample('D').sum()
```

| Date       | Units |
|------------|-------|
| 2015-02-02 | 6.0   |
| 2015-02-03 | 13.0  |
| 2015-02-04 | 13.5  |
| 2015-02-05 | 14.5  |
| 2015-02-06 | NaN   |
| 2015-02-07 | 1.0   |
| 2015-02-08 | NaN   |
| 2015-02-09 | 13.0  |
| 2015-02-10 | NaN   |
| 2015-02-11 | 5.5   |
| 2015-02-12 | NaN   |
| 2015-02-13 | NaN   |

# Method chaining

```
sales.resample('D').sum().max()
```

```
Units    29.0  
dtype: float64
```

# Resampling strings

```
sales.resample('W').count()
```

|            | Company | Product | Units |
|------------|---------|---------|-------|
| Date       |         |         |       |
| 2015-02-08 | 8       | 8       | 8     |
| 2015-02-15 | 4       | 4       | 4     |
| 2015-02-22 | 5       | 5       | 5     |
| 2015-03-01 | 2       | 2       | 2     |

# Resampling frequencies

| Input      | Description  |
|------------|--------------|
| 'min', 'T' | minute       |
| 'H'        | hour         |
| 'D'        | day          |
| 'B'        | business day |
| 'W'        | week         |
| 'M'        | month        |
| 'Q'        | quarter      |
| 'A'        | year         |



# Multiplying frequencies

```
sales.loc[:, 'Units'].resample('2W').sum()
```

```
Date
2015-02-08    82
2015-02-22    79
2015-03-08    14
Freq: 2W-SUN, Name: Units, dtype: int64
```

# Upsampling

```
two_days = sales.loc['2015-2-4': '2015-2-5', 'Units']  
two_days
```

```
Date  
2015-02-04 15:30:00    13  
2015-02-04 22:00:00    14  
2015-02-05 02:00:00    19  
2015-02-05 22:00:00    10  
Name: Units, dtype: int64
```

# Upsampling and filling

```
two_days.resample('4H').ffill()
```

```
Date
Date
2015-02-04 12:00:00      NaN
2015-02-04 16:00:00    13.0
2015-02-04 20:00:00    13.0
2015-02-05 00:00:00    14.0
2015-02-05 04:00:00    19.0
2015-02-05 08:00:00    19.0
2015-02-05 12:00:00    19.0
2015-02-05 16:00:00    19.0
2015-02-05 20:00:00    19.0
Freq: 4H, Name: Units, dtype: float64
```

# Let's practice!

PANDAS FOUNDATIONS

# Manipulating time series data

PANDAS FOUNDATIONS



**Dhavide Aruliah**

Director of Training, Anaconda

# Sales data

```
import pandas as pd
sales = pd.read_csv('sales-feb-2015.csv',
                    parse_dates=['Date'])
sales.head()
```

|   | Date                | Company         | Product  | Units |
|---|---------------------|-----------------|----------|-------|
| 0 | 2015-02-02 08:30:00 | Hooli           | Software | 3     |
| 1 | 2015-02-02 21:00:00 | Mediacore       | Hardware | 9     |
| 2 | 2015-02-03 14:00:00 | Initech         | Software | 13    |
| 3 | 2015-02-04 15:30:00 | Streeplex       | Software | 13    |
| 4 | 2015-02-04 22:00:00 | Acme Coporation | Hardware | 14    |

# String methods

```
sales['Company'].str.upper()
```

```
0      HOOLI
1    MEDIACORE
2    INITECH
3    STREEPLEX
4  ACME COPORATION
5  ACME COPORATION
6      HOOLI
7  ACME COPORATION
8    STREEPLEX
9    MEDIACORE
10     INITECH
11     HOOLI
12     HOOLI
13    MEDIACORE
14    MEDIACORE
15    MEDIACORE
...
```

# Substring matching

```
sales['Product'].str.contains('ware')
```

```
0      True
1      True
2      True
3      True
4      True
5      True
6     False
7      True
8     False
9      True
10     True
11     True
12     True
13     True
14     False
...
```



# Boolean arithmetic

`True + False`

1

`True + True`

2

`False + False`

0

# Boolean reduction

```
sales[ 'Product' ].str.contains( 'ware' ).sum()
```

```
14
```

# Datetime methods

```
sales['Date'].dt.hour
```

```
0      8
1     21
2     14
3     15
4     22
5      2
6     22
7     23
8      9
9     13
10    20
11    23
12    12
13    11
14    16
...
```

# Set timezone

```
central = sales['Date'].dt.tz_localize('US/Central')
central
```

```
0    2015-02-02 08:30:00-06:00
1    2015-02-02 21:00:00-06:00
2    2015-02-03 14:00:00-06:00
3    2015-02-04 15:30:00-06:00
4    2015-02-04 22:00:00-06:00
5    2015-02-05 02:00:00-06:00
6    2015-02-05 22:00:00-06:00
7    2015-02-07 23:00:00-06:00
8    2015-02-09 09:00:00-06:00
9    2015-02-09 13:00:00-06:00
10   2015-02-11 20:00:00-06:00
11   2015-02-11 23:00:00-06:00
12   2015-02-16 12:00:00-06:00
```

...

```
Name: Date, dtype: datetime64[ns, US/Central]
```

# Convert timezone

```
central.dt.tz_convert('US/Eastern')
```

```
0    2015-02-02 09:30:00-05:00
1    2015-02-02 22:00:00-05:00
2    2015-02-03 15:00:00-05:00
3    2015-02-04 16:30:00-05:00
4    2015-02-04 23:00:00-05:00
5    2015-02-05 03:00:00-05:00
6    2015-02-05 23:00:00-05:00
7    2015-02-08 00:00:00-05:00
8    2015-02-09 10:00:00-05:00
9    2015-02-09 14:00:00-05:00
10   2015-02-11 21:00:00-05:00
11   2015-02-12 00:00:00-05:00
12   2015-02-16 13:00:00-05:00
13   2015-02-19 12:00:00-05:00
14   2015-02-19 17:00:00-05:00
```

...

```
Name: Date, dtype: datetime64[ns, US/Eastern]
```

# Method chaining

```
sales['Date'].dt.tz_localize('US/Central').  
    dt.tz_convert('US/Eastern')
```

```
0    2015-02-02 09:30:00-05:00  
1    2015-02-02 22:00:00-05:00  
2    2015-02-03 15:00:00-05:00  
3    2015-02-04 16:30:00-05:00  
4    2015-02-04 23:00:00-05:00  
5    2015-02-05 03:00:00-05:00  
6    2015-02-05 23:00:00-05:00  
7    2015-02-08 00:00:00-05:00  
8    2015-02-09 10:00:00-05:00  
9    2015-02-09 14:00:00-05:00  
10   2015-02-11 21:00:00-05:00  
11   2015-02-12 00:00:00-05:00  
12   2015-02-16 13:00:00-05:00  
13   2015-02-19 12:00:00-05:00  
14   2015-02-19 17:00:00-05:00  
...  
Name: Date, dtype: datetime64[ns, US/Eastern]
```

# World Population

```
population = pd.read_csv('world_population.csv',  
    parse_dates=True, index_col= 'Date')  
population
```

|            | Population   |
|------------|--------------|
| Date       |              |
| 1960-12-31 | 2.087485e+10 |
| 1970-12-31 | 2.536513e+10 |
| 1980-12-31 | 3.057186e+10 |
| 1990-12-31 | 3.644928e+10 |
| 2000-12-31 | 4.228550e+10 |
| 2010-12-31 | 4.802217e+10 |

# Upsample population

```
population.resample('A').first()
```

| Date       | Population   |
|------------|--------------|
| 1960-12-31 | 2.087485e+10 |
| 1961-12-31 | NaN          |
| 1962-12-31 | NaN          |
| 1963-12-31 | NaN          |
| 1964-12-31 | NaN          |
| 1965-12-31 | NaN          |
| 1966-12-31 | NaN          |
| 1967-12-31 | NaN          |
| 1968-12-31 | NaN          |
| 1969-12-31 | NaN          |
| 1970-12-31 | 2.536513e+10 |
| 1971-12-31 | NaN          |
| 1972-12-31 | NaN          |



# Interpolate missing data

```
population.resample('A').first().interpolate('linear')
```

| Date       | Population   |
|------------|--------------|
| 1960-12-31 | 2.087485e+10 |
| 1961-12-31 | 2.132388e+10 |
| 1962-12-31 | 2.177290e+10 |
| 1963-12-31 | 2.222193e+10 |
| 1964-12-31 | 2.267096e+10 |
| 1965-12-31 | 2.311999e+10 |
| 1966-12-31 | 2.356902e+10 |
| 1967-12-31 | 2.401805e+10 |
| 1968-12-31 | 2.446707e+10 |
| 1969-12-31 | 2.491610e+10 |
| 1970-12-31 | 2.536513e+10 |
| 1971-12-31 | 2.588580e+10 |
| 1972-12-31 | 2.640648e+10 |

# Let's practice!

PANDAS FOUNDATIONS

# Time series visualization

PANDAS FOUNDATIONS



**Dhavide Aruliah**

Director of Training, Anaconda

# Topics

- Line types
- Plot types
- Subplots

# S&P 500 Data

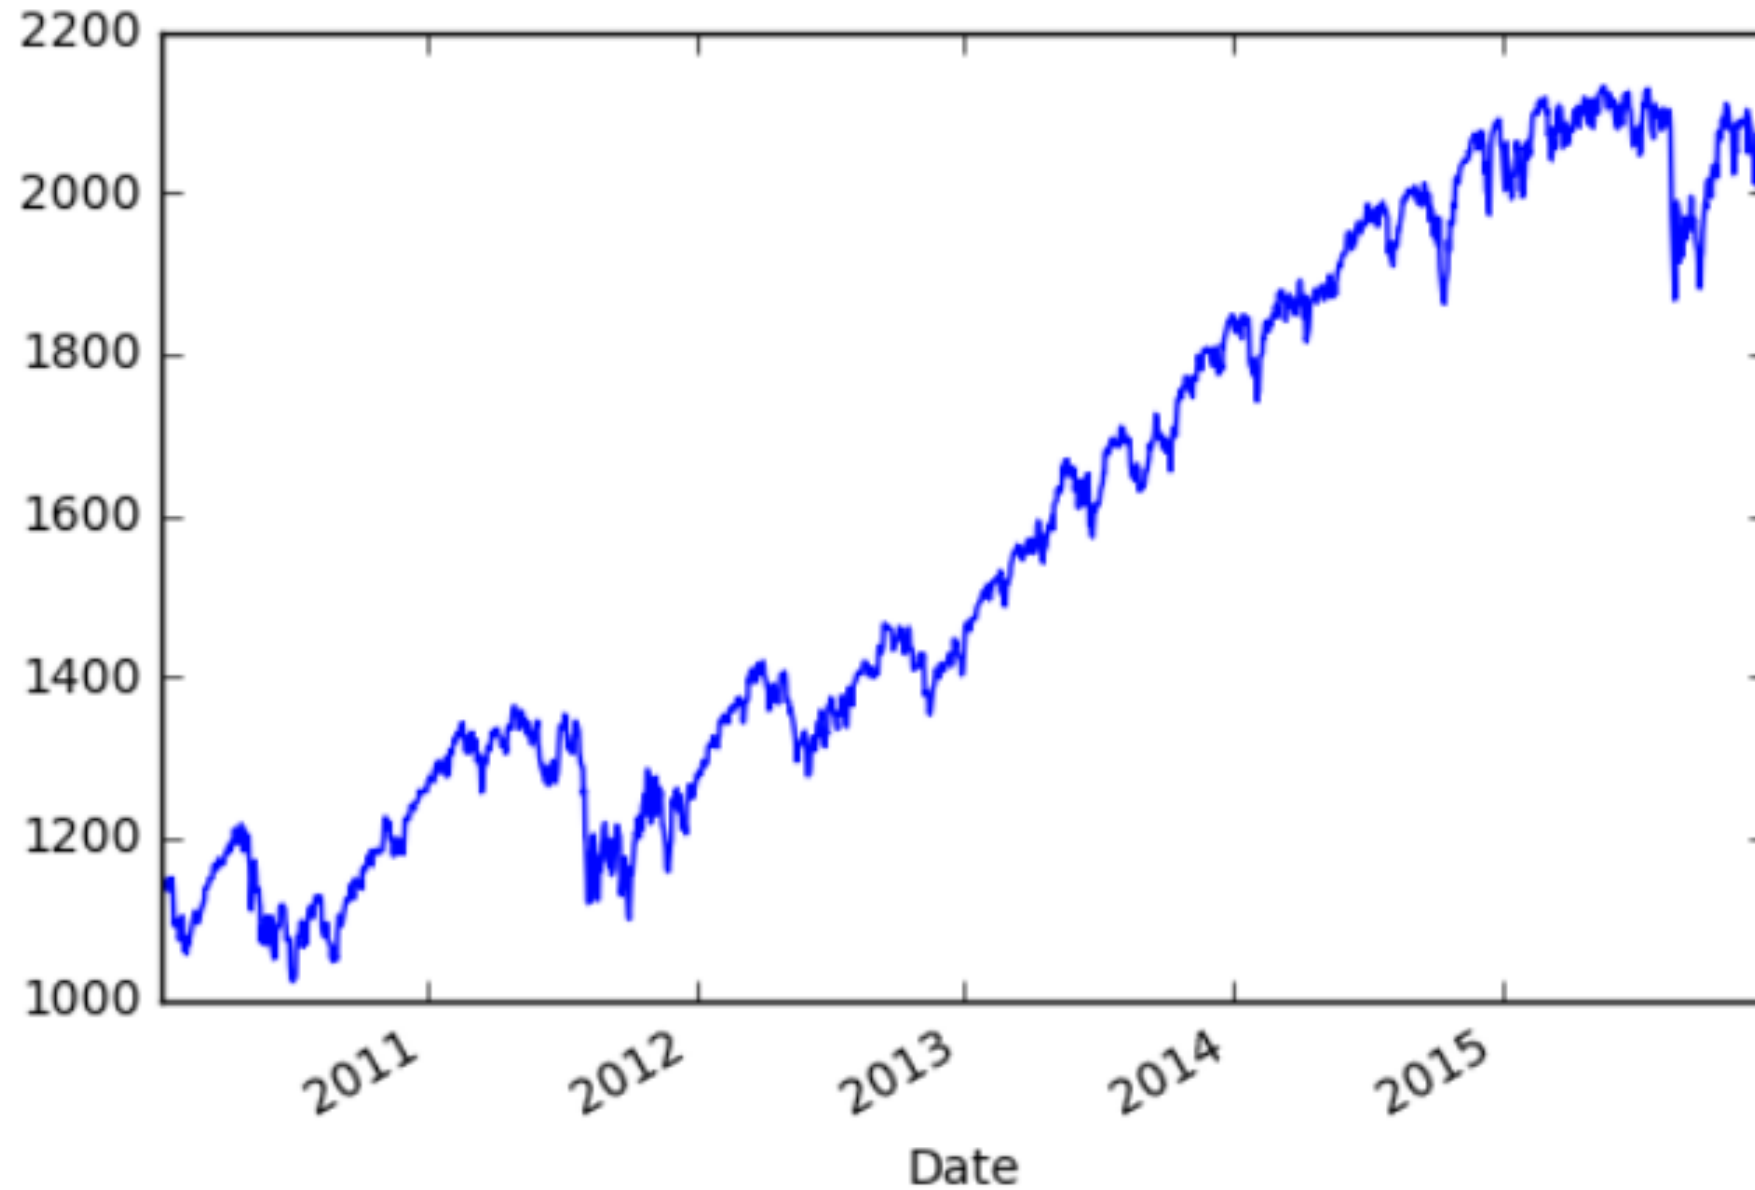
```
import pandas as pd
import matplotlib.pyplot as plt
sp500 = pd.read_csv('sp500.csv', parse_dates=True,
                    index_col= 'Date')
sp500.head()
```

|            | Open        | High        | Low         | Close       | Volume     | Adj Close   |
|------------|-------------|-------------|-------------|-------------|------------|-------------|
| Date       |             |             |             |             |            |             |
| 2010-01-04 | 1116.560059 | 1133.869995 | 1116.560059 | 1132.989990 | 3991400000 | 1132.989990 |
| 2010-01-05 | 1132.660034 | 1136.630005 | 1129.660034 | 1136.520020 | 2491020000 | 1136.520020 |
| 2010-01-06 | 1135.709961 | 1139.189941 | 1133.949951 | 1137.140015 | 4972660000 | 1137.140015 |
| 2010-01-07 | 1136.270020 | 1142.459961 | 1131.319946 | 1141.689941 | 5270680000 | 1141.689941 |
| 2010-01-08 | 1140.520020 | 1145.390015 | 1136.219971 | 1144.979980 | 4389590000 | 1144.979980 |

# Pandas plot

```
sp500[ 'Close' ].plot()  
plt.show()
```

# Default plot

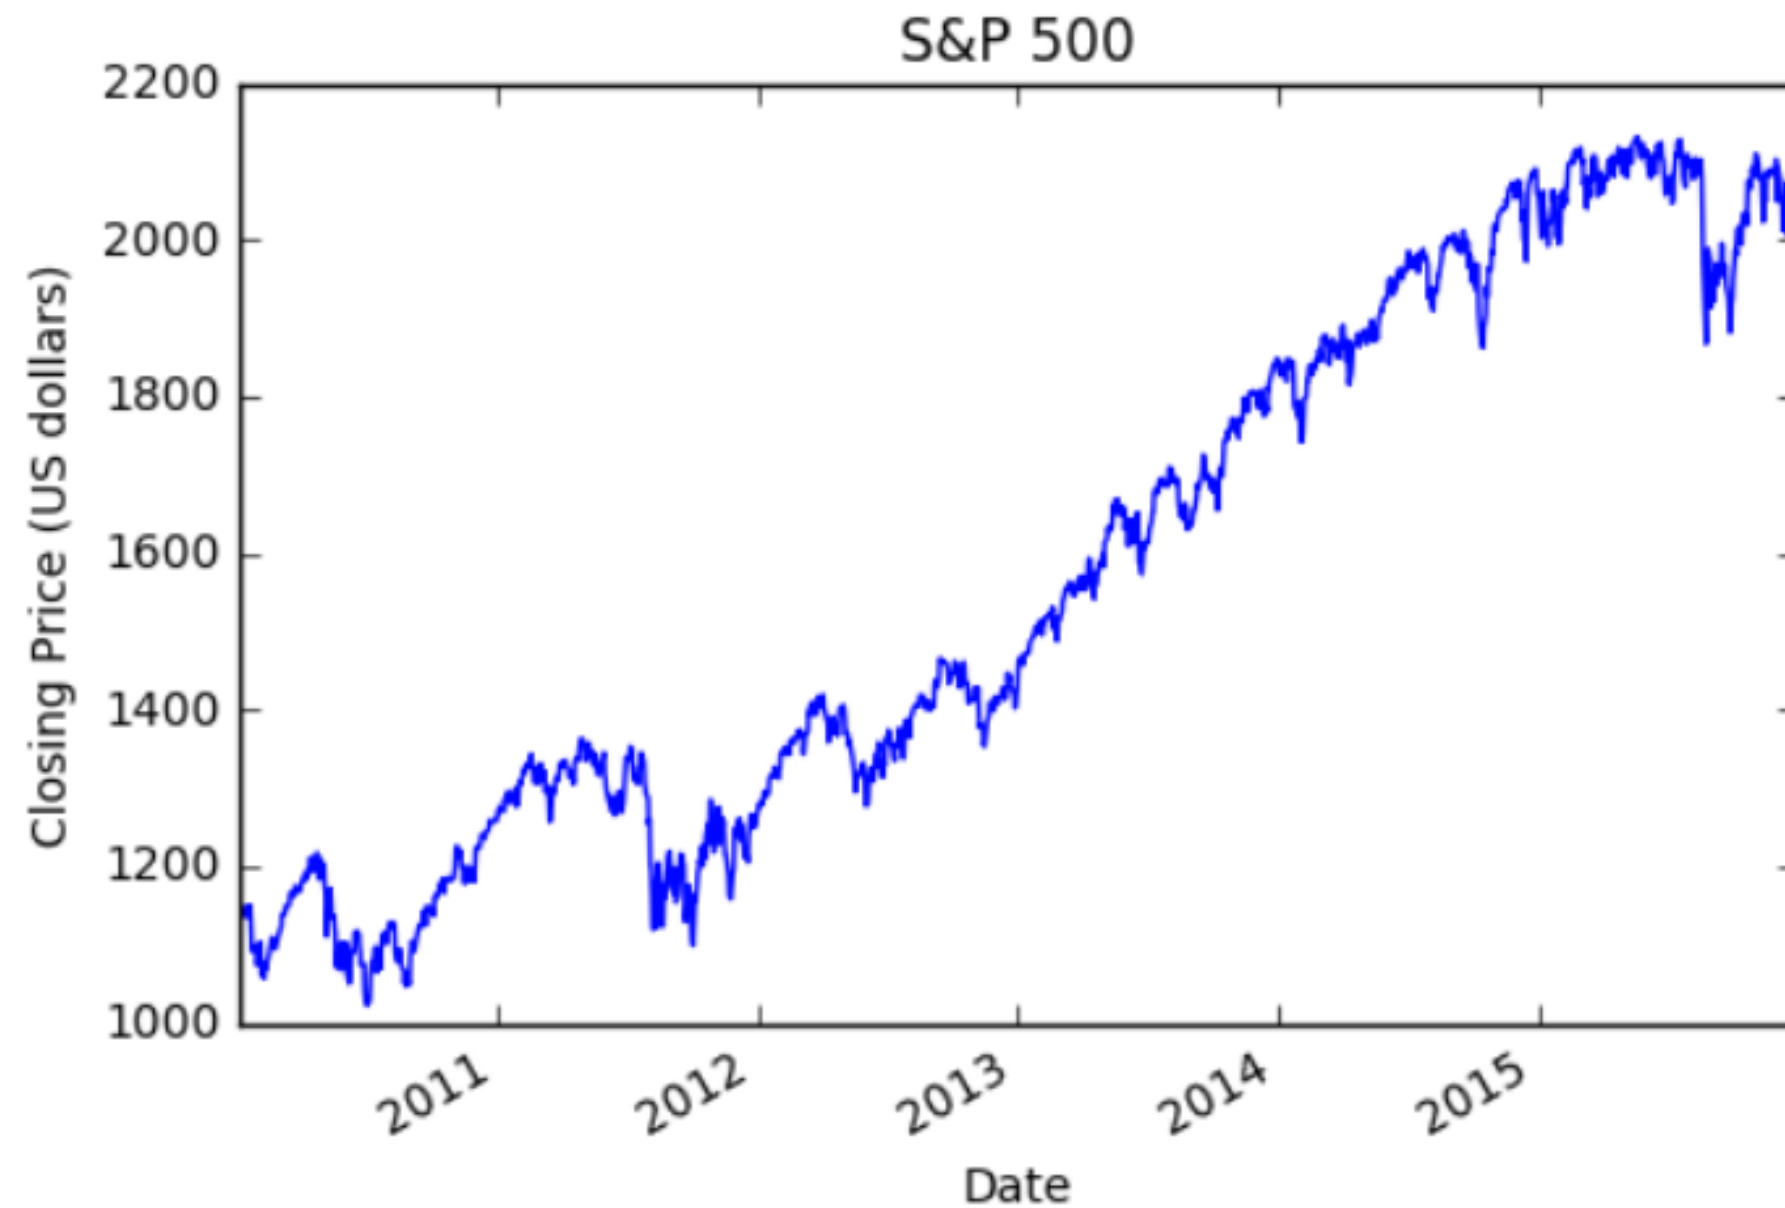


# Labels and title

```
sp500['Close'].plot(title='S&P 500')  
plt.ylabel('Closing Price (US Dollars)')  
plt.show()
```



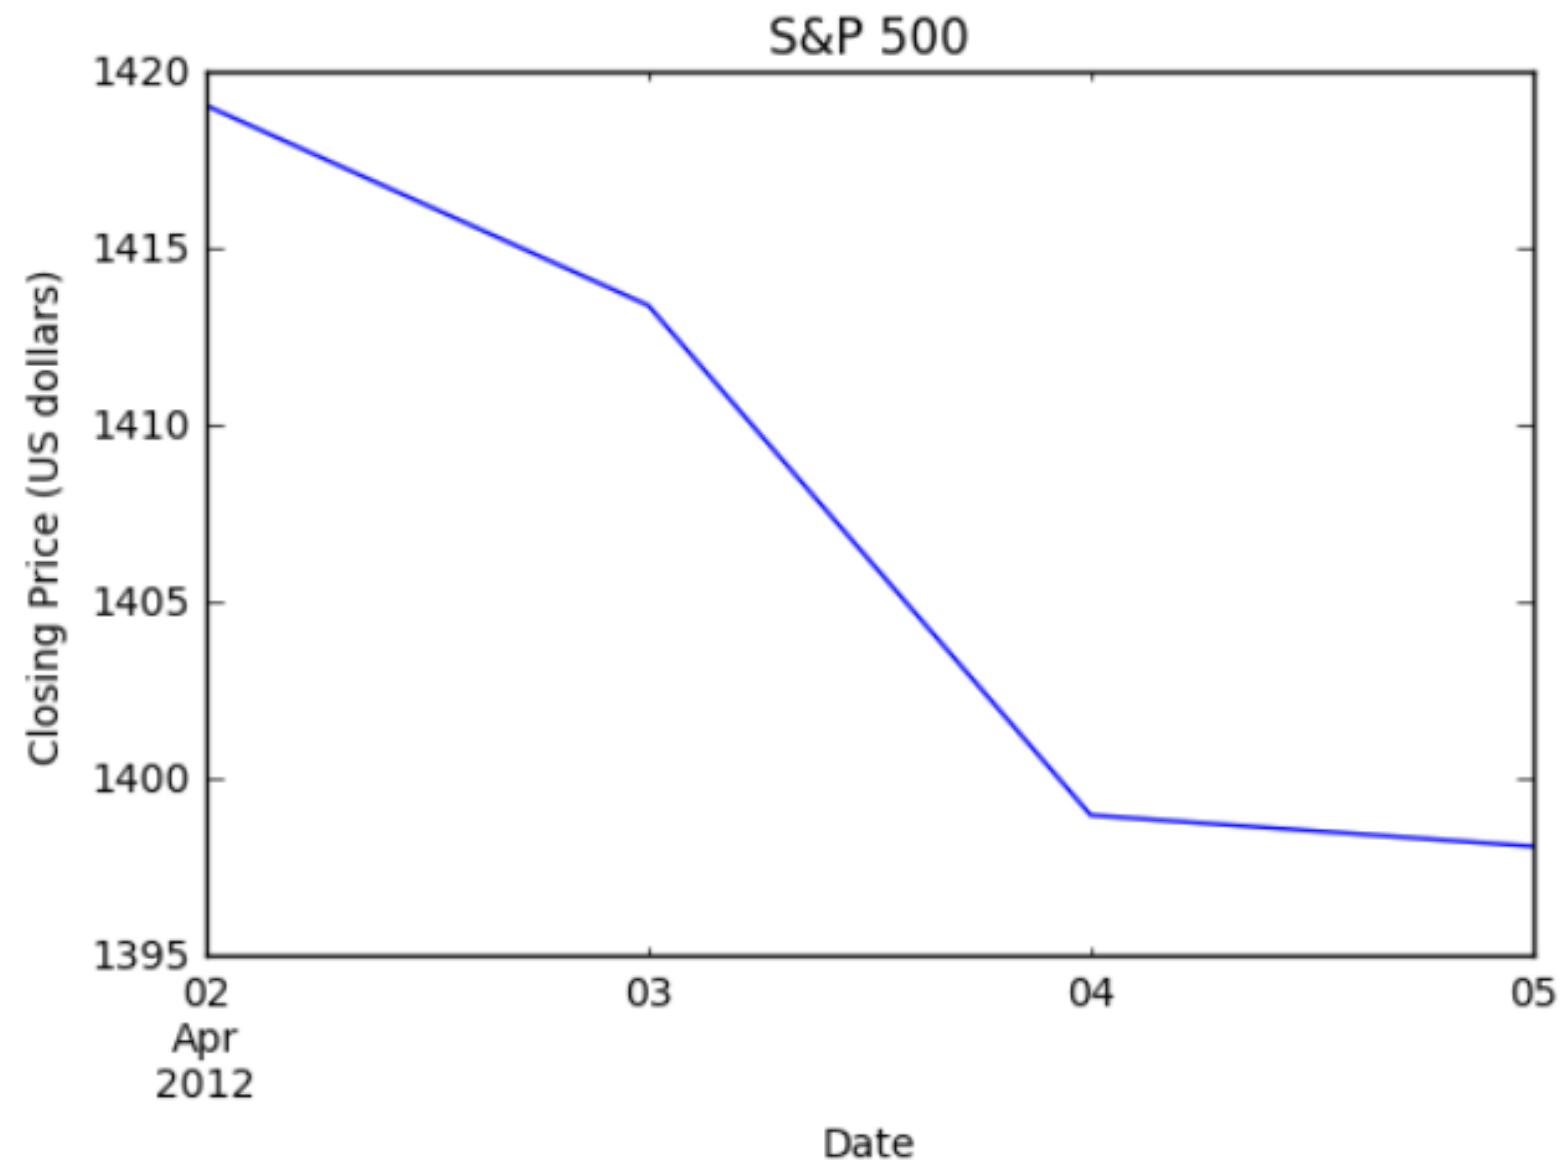
# Labels and title



# One week

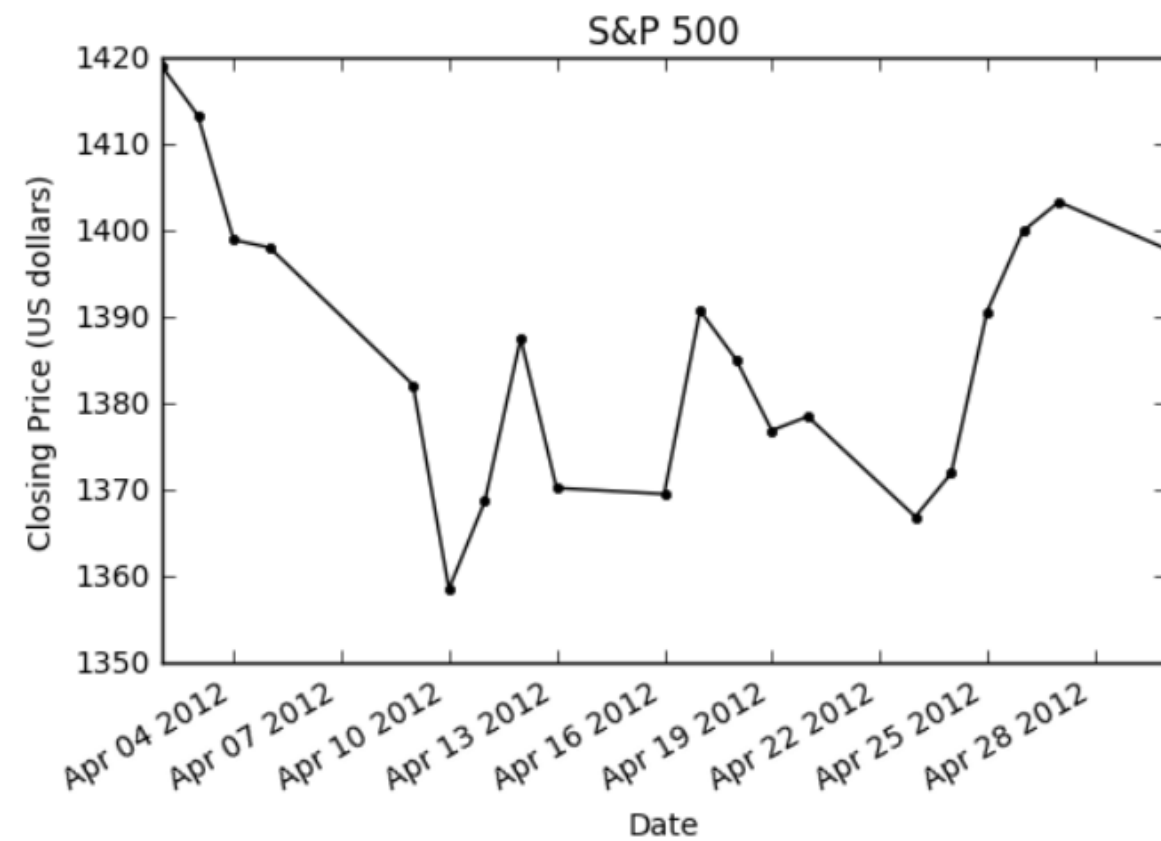
```
sp500.loc['2012-4-1':'2012-4-7', 'Close'].plot(title='S&P  
500')  
plt.ylabel('Closing Price (US Dollars)')  
plt.show()
```

# One week



# Plot styles

```
sp500.loc['2012-4', 'Close'].plot(style='k.-',  
                                  title='S&P500')  
plt.ylabel('Closing Price (US Dollars)')  
plt.show()
```



# More plot styles


- Style format string
  - color (k: black)
  - marker (.: dot)
  - line type (-: solid)

# More plot styles

| Color    | Marker    | Line      |
|----------|-----------|-----------|
| b: blue  | o: circle | : dotted  |
| g: green | *: star   | —: dashed |
| r: red   | s: square |           |
| c: cyan  | +: plus   |           |

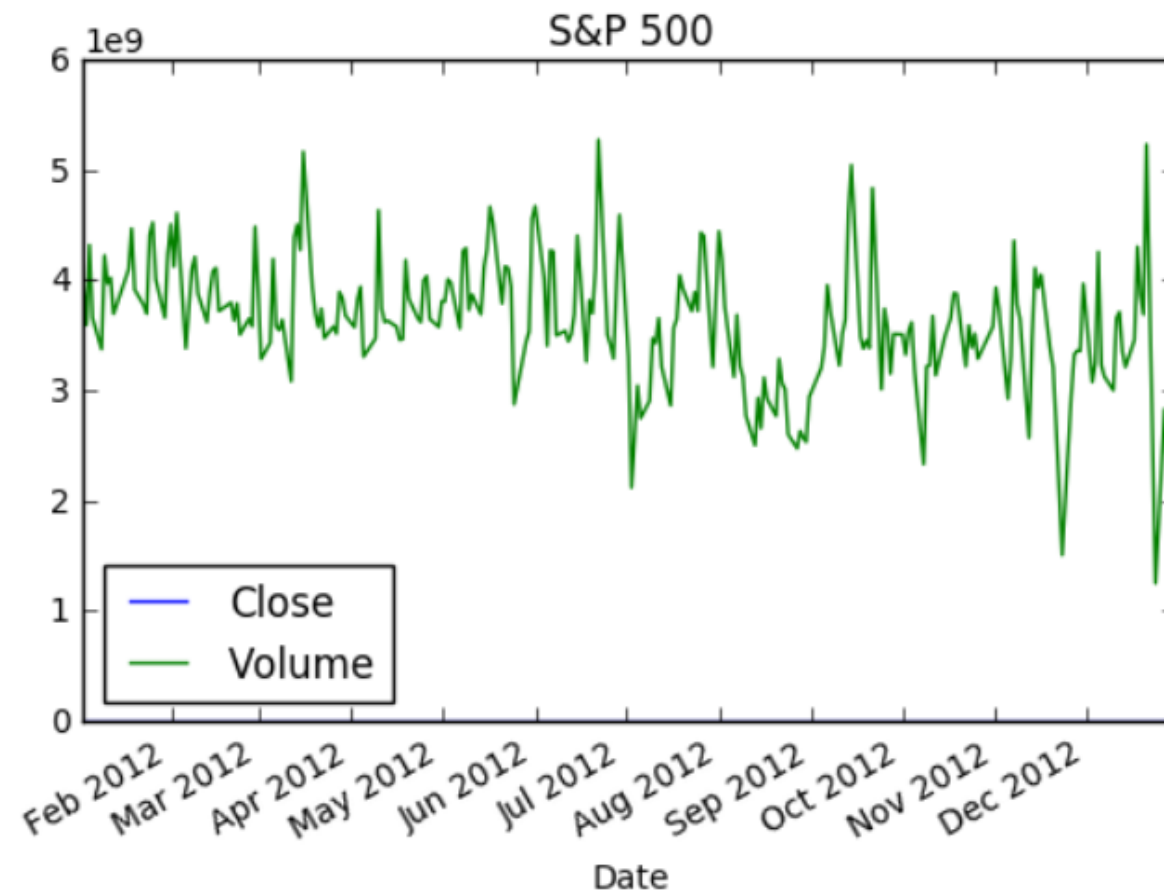
# Area plot

```
sp500['Close'].plot(kind='area', title='S&P 500')  
plt.ylabel('Closing Price (US Dollars)')  
plt.show()
```



# Multiple columns

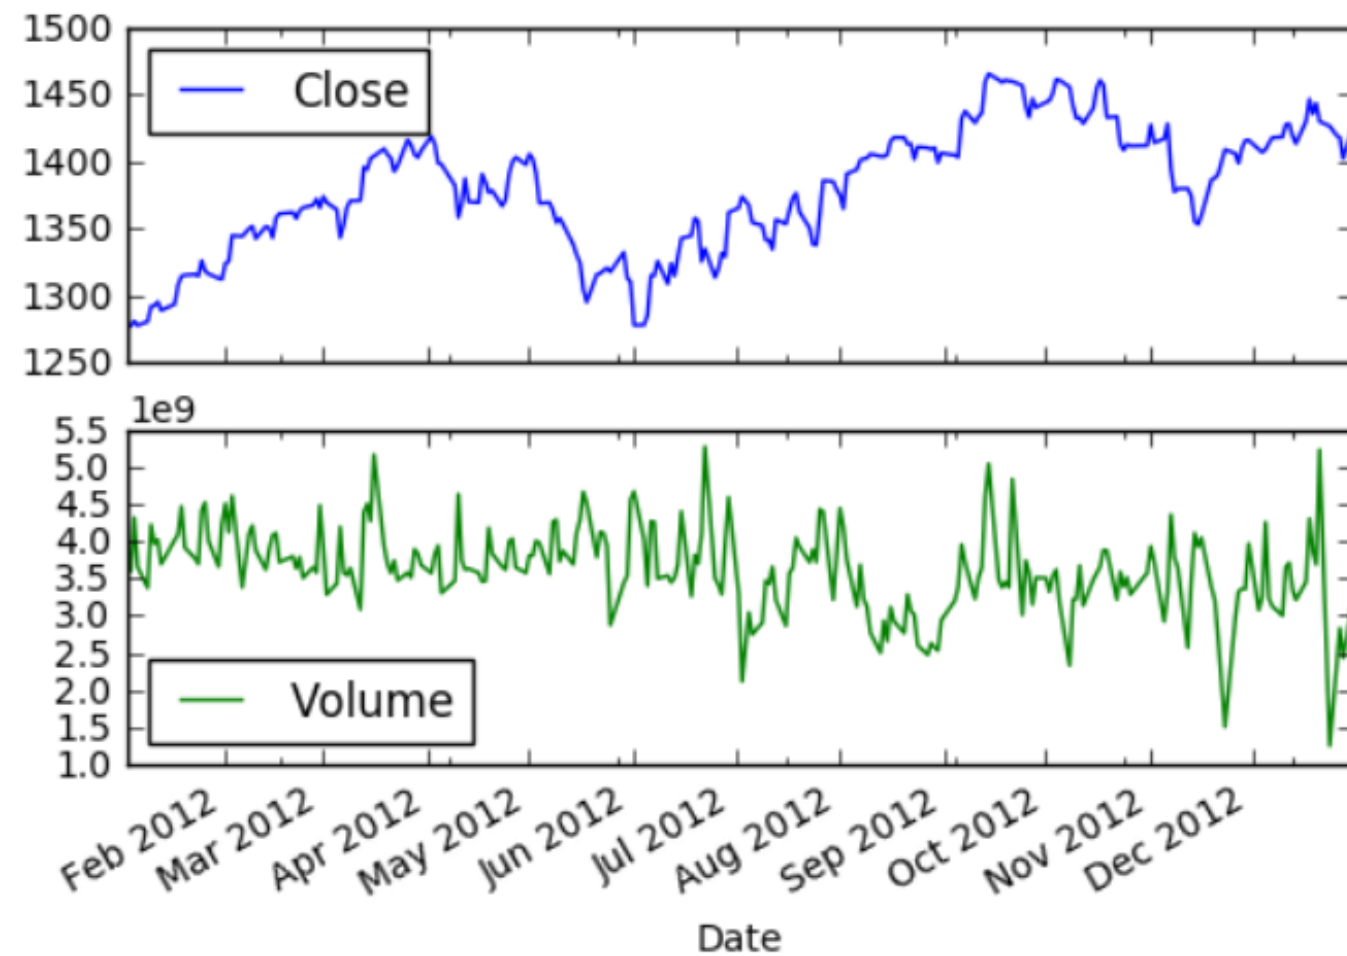
```
sp500.loc['2012', ['Close', 'Volume']].plot(title='S&P  
500')  
plt.show()
```





# Subplots

```
sp500.loc['2012', ['Close', 'Volume']].plot(subplots=True)  
plt.show()
```



# Let's practice!

PANDAS FOUNDATIONS